

Latency Scaling with Token Count

